


Token-Efficient Multimodal Reasoning via Scene-to-Structured-Text (SST) Compression



Current SST branch is an oracle study because SST is derived from GQA scene graphs rather than a fully automatic image-to-SST extractor.